

Indexes Best Practices (II)

More T-SQL

Control-Of-Flow Language

Indexes Best Practices (II)

Indexed Views

SET options	Required value	Default server value
ANSI_NULLS	ON	ON
ANSI_PADDING	ON	ON
ANSI_WARNINGS	ON	ON
ARITHABORT	ON	ON
CONCAT_NULL_YIELDS_NULL	ON	ON
NUMERIC_ROUNDABORT	OFF	OFF
QUOTED_IDENTIFIER	ON	ON

Indexed Views Restrictions

- SELECT statement cannot reference other views
- All functions must be deterministic
- AVG, MIN, MAX, STDEV, STDEVP, VAR and VARP are not allowed
- The index must be both clustered and unique
- SELECT statement must not contain subqueries, outer joins, EXCEPT, INTERSECT, TOP, UNION, ORDER BY, DISTINCT etc

Columnstore Indexes

- Groups and stores data for each column and then joins all the columns to complete the whole index
- Suited for warehouses (read only tables)
- Up to 10x query performance (vs. traditional row-oriented storage)
- Up to 10x data compression over the uncompressed data size
- The same table can have both row store index and column store index, The *Query Optimizer* will decide when to use the column store index and when to use other types of indexes

Row store for B-Tree or Heap

Row 1	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 2	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 3	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 4	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 5	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10

Page 1

Row 6	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 7	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 8	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
.....	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row n	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10

Page 2

Column Store Index

Row 1	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 2	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 3	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 4	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 5	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 6	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 7	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row 8	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
.....	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
Row n	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
	Page 1	Page 2	Page 3	Page 4	Page 5	Page 6	Page 7	Page 8	Page 9	Page 10

Hard and fast rules for indexing

- Each table should have a clustered index that is (ideally) small, selective, ever increasing, and static. (a table without a clustered index is called a *heap*.)
- Implement nonclustered indexes on foreign key relationships
- Implement nonclustered indexes on columns that are frequently used in WHERE clauses.
- Do not implement single-column indexes on every column in a table. This will cause high overhead.
- In multi-column indexes, list the most selective (nearest to unique) first in the column list.
- For most often-used queries create covering nonclustered index.

Fragmentation

- *Internal Fragmentation*: records are stored non-contiguously inside the page. Internal fragmentation occurs if there is unused space between records in a page. The fullness of each page can vary over time. This unused space causes poor cache utilization and more I/O, which ultimately leads to poor query performance.

Fragmentation

- *External Fragmentation:* When on disk, the physical storage of pages and extents is not contiguous. When the extents of a table are not physically stored contiguously on disk, switching from one extent to another causes higher disk rotations.

Fragmentation

- *Logical Fragmentation*: Every index page is linked with previous and next page in the logical order of column data. Because of Page Split, the pages turn into *out-of-order* pages.
- An *out-of-order* page is a page for which the next physical page allocated to the index is not the page pointed to by the next-page pointer in the current leaf page.


Page read requests: 2


Extent switches: 0

Disk space used by table: 16 KB

avg_fragmentation_in_percent: 0

avg_page_space_used_in_percent: 100





Page read requests: 6

Extent switches: 5

Disk space used by table: 48
KB

avg_fragmentation_in_
percent > 80

avg_page_space_used_in_p
ercent: 33

Fragmentation

- *sys.dm_db_index_physical_stats*
 - **avg_fragmentation_in_percent:** This is a percentage value that represents external fragmentation.
 - **avg_page_space_used_in_percent:** This is an average percentage use of pages that represents to internal fragmentation.
- **Reducing Fragmentation in a Heap:**
 - To reduce the fragmentation of a heap, create a clustered index on the table.
 - Creating the clustered index: rearrange the records in an order, and then place the pages contiguously on disk.

Fragmentation

Reducing Fragmentation in a Index:

- If `avg_fragmentation_in_percent` $> 5\%$ and $< 30\%$, then use `ALTER INDEX REORGANIZE`:
 - reorder the leaf level pages of the index in a logical order.
- If `avg_fragmentation_in_percent` $> 30\%$, then use `ALTER INDEX REBUILD`:
 - replacement for `DBCC DBREINDEX` to rebuild the index online or offline. In such case, we can also use the drop and re-create index method.
- Drop and re-create the clustered index:
 - Re-creating a clustered index redistributes the data and results in full data pages. The level of fullness can be configured by using the `FILLFACTOR` option in `CREATE INDEX`.

More T-SQL

Control-Of-Flow Language

Control-of-Flow Language

BEGIN...END

RETURN

BREAK

THROW

CONTINUE

TRY...CATCH

GOTO label

WAITFOR

IF...ELSE

WHILE

RETURN

RETURN [integer_expression]

- exits unconditionally from a query or procedure
- returning from a procedure
- returning status codes
 - stored procs return 0 (success), or
 - a nonzero value (failure)

RETURN

```
CREATE PROCEDURE checkstate @param varchar(11)
```

```
AS
```

```
    IF @param= 'WA'
```

```
        RETURN 1
```

```
    ELSE
```

```
        RETURN 2;
```

```
GO
```

```
-----
```

```
DECLARE @return_status int;
```

```
EXEC @return_status = checkstate 'AK';
```

```
GO
```

WHILE

WHILE Boolean_expression

{ sql_statement | statement_block | BREAK |
CONTINUE }

- sets a condition for the repeated execution of an SQL statement or statement block

BREAK

- exits the innermost loop in a WHILE statement or an IF...ELSE statement inside a WHILE loop.

CONTINUE

- restarts a WHILE loop; any statements after the CONTINUE keyword are ignored

GOTO

- alters the flow of execution to a label

Label:

GOTO Label

WAITFOR

```
WAITFOR { DELAY 'time_to_pass' |  
          TIME 'time_to_execute' |  
          [ ( receive_statement ) |  
            ( get_conversation_group_statement ) ]  
          [ , TIMEOUT timeout ] }
```

- blocks the execution of a batch, stored procedure, or transaction

WAITFOR

- execution continues at 08:35

WAITFOR TIME '08:35';

- execution continues after 2 hours

WAITFOR DELAY '02:00';

- if the server is busy → the counter does not start immediately → the delay may be longer than specified.

THROW

```
THROW [  
    { error_number | @local_variable },  
    { message | @local_variable },  
    { state | @local_variable } ] [ ; ]
```

- raises an exception and transfers execution to a CATCH block of a TRY...CATCH construct
- the exception severity is always set to 16.

```
THROW 51000, 'Record does not exist', 1;
```

TRY ... CATCH

BEGIN TRY

 { sql_statement | statement_block }

END TRY

BEGIN CATCH

 [{ sql_statement | statement_block }]

END CATCH [;]

- implements error handling for Transact-SQL
- catches all execution errors that have a severity >10 that do not close the database connection

TRY ... CATCH

- `ERROR_NUMBER()` returns the error number
- `ERROR_SEVERITY()` returns the severity
- `ERROR_STATE()` returns the error state number
- `ERROR_PROCEDURE()` returns the name of the stored procedure/trigger where the error occurred
- `ERROR_LINE()` returns the line number that caused the error
- `ERROR_MESSAGE()` returns the error message

Error messages

- error number
 - integer value between 1 and 49999
 - custom error messages: 50001...
- error severity
 - 26 severity levels
 - error with severity level ≥ 16 are logged automatically
 - error with severity level between 20 and 25 are fatal and the connection is terminated
- error message: up to 255 chars